



SK Big Data Course (2019)

# 빅데이터 아키텍처

SK Big Data Course (2019), 빅데이터 아키텍처

# Designing Modern Streaming Data Applications



# Me:

About Me

- **김영우 (Youngwoo Kim)**
- **소프트웨어 엔지니어**
- **Hi-Tech DT 팀, Data Labs, ICT기술센터, SK 텔레콤**

- 개요
- 빅데이터 아키텍처
  - 스트리밍 데이터 응용을 위한 빅데이터 아키텍처
    - a.k.a, Fast data
    - a.k.a, Real-time data
- 스트리밍 데이터 플랫폼
  - Data Source
  - Event Hub / Message Broker
  - Data Ingestion / Data Integration
  - Data Storage
  - Stream Processing
  - Data Analytics / SQL / Search
- **Hands-on Labs**

# 개요



- **스트리밍 데이터 처리 시스템**을 구성하기 위해 필요한 구성 요소에 대하여 학습
- 실시간 데이터를 위한 **오픈소스** 프로젝트와 해당 프로젝트의 특징 학습
- 성공적인 데이터 응용 설계와 개발을 위한 고려사항
- 스트리밍 데이터 아키텍처 시나리오를 바탕으로 **e2e 스트리밍 응용 개발** 실습

# 빅데이터?

- nVs

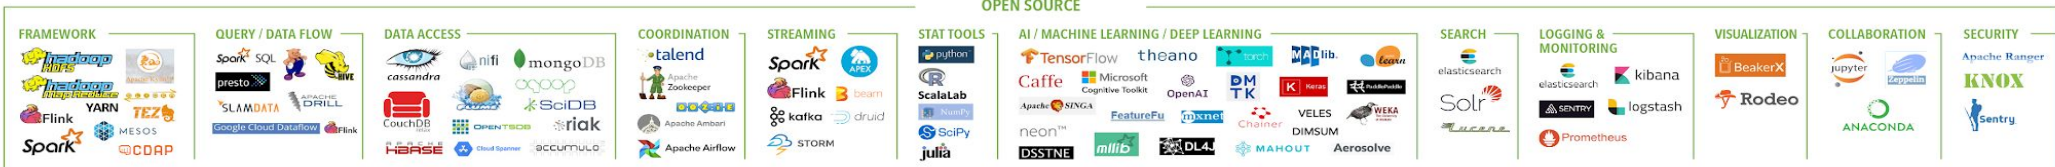
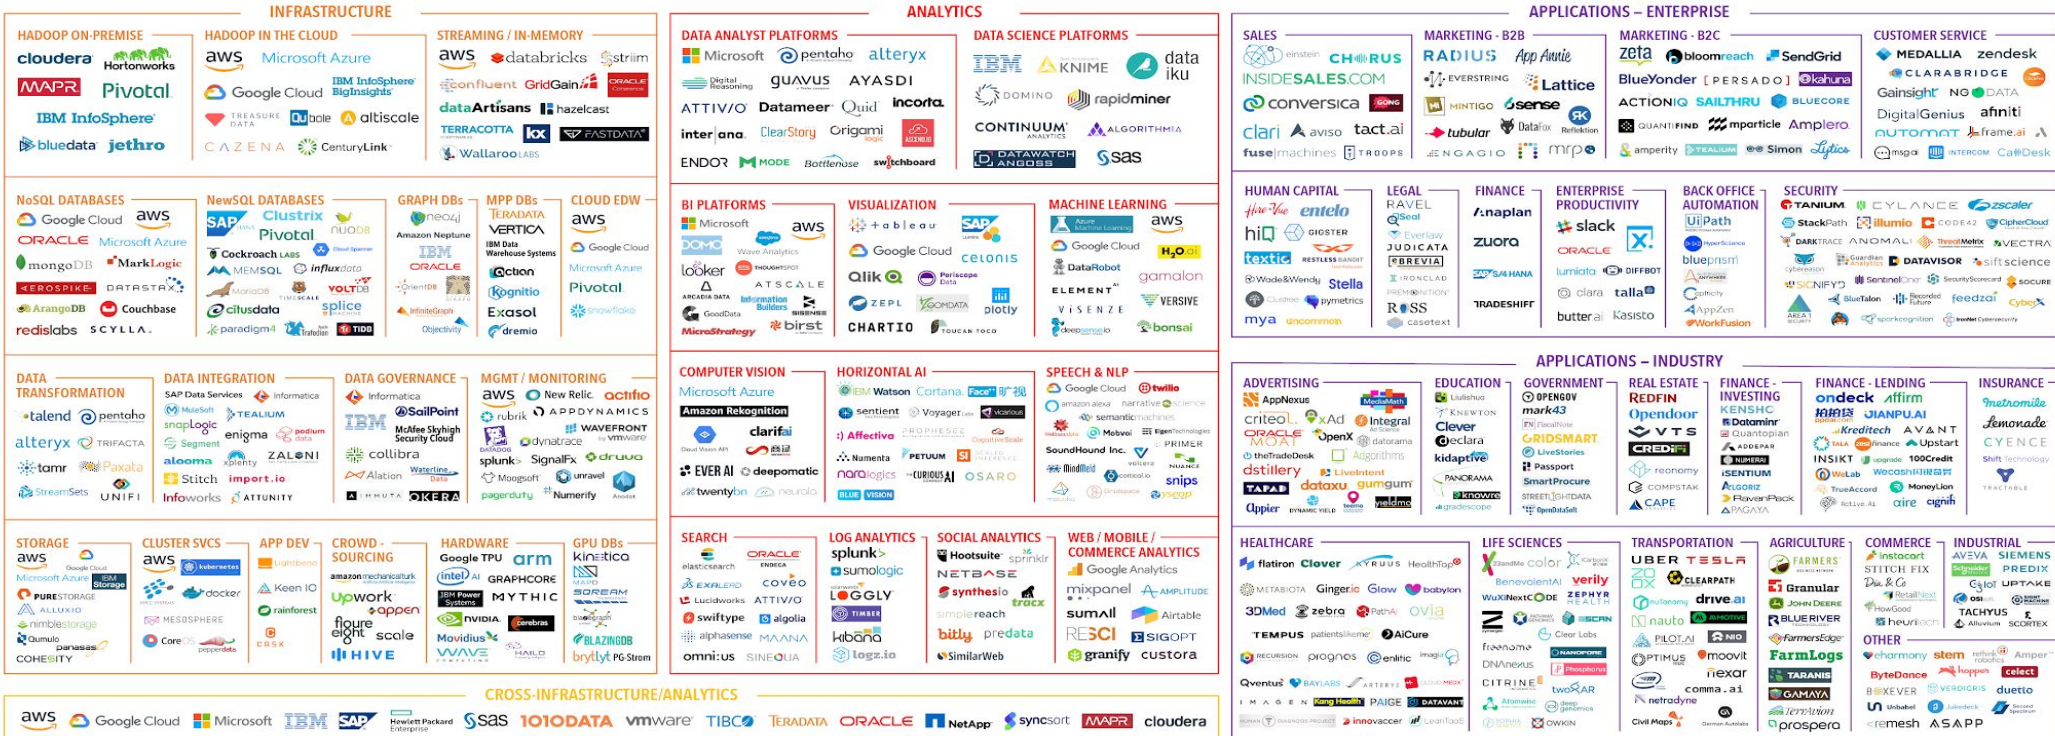
- 3Vs
- 4Vs
- 5Vs
- 6Vs
- 7Vs
- .....
- As of May 2019
  - 8Vs
  - 10Vs
  - ㅋㅋㅋ





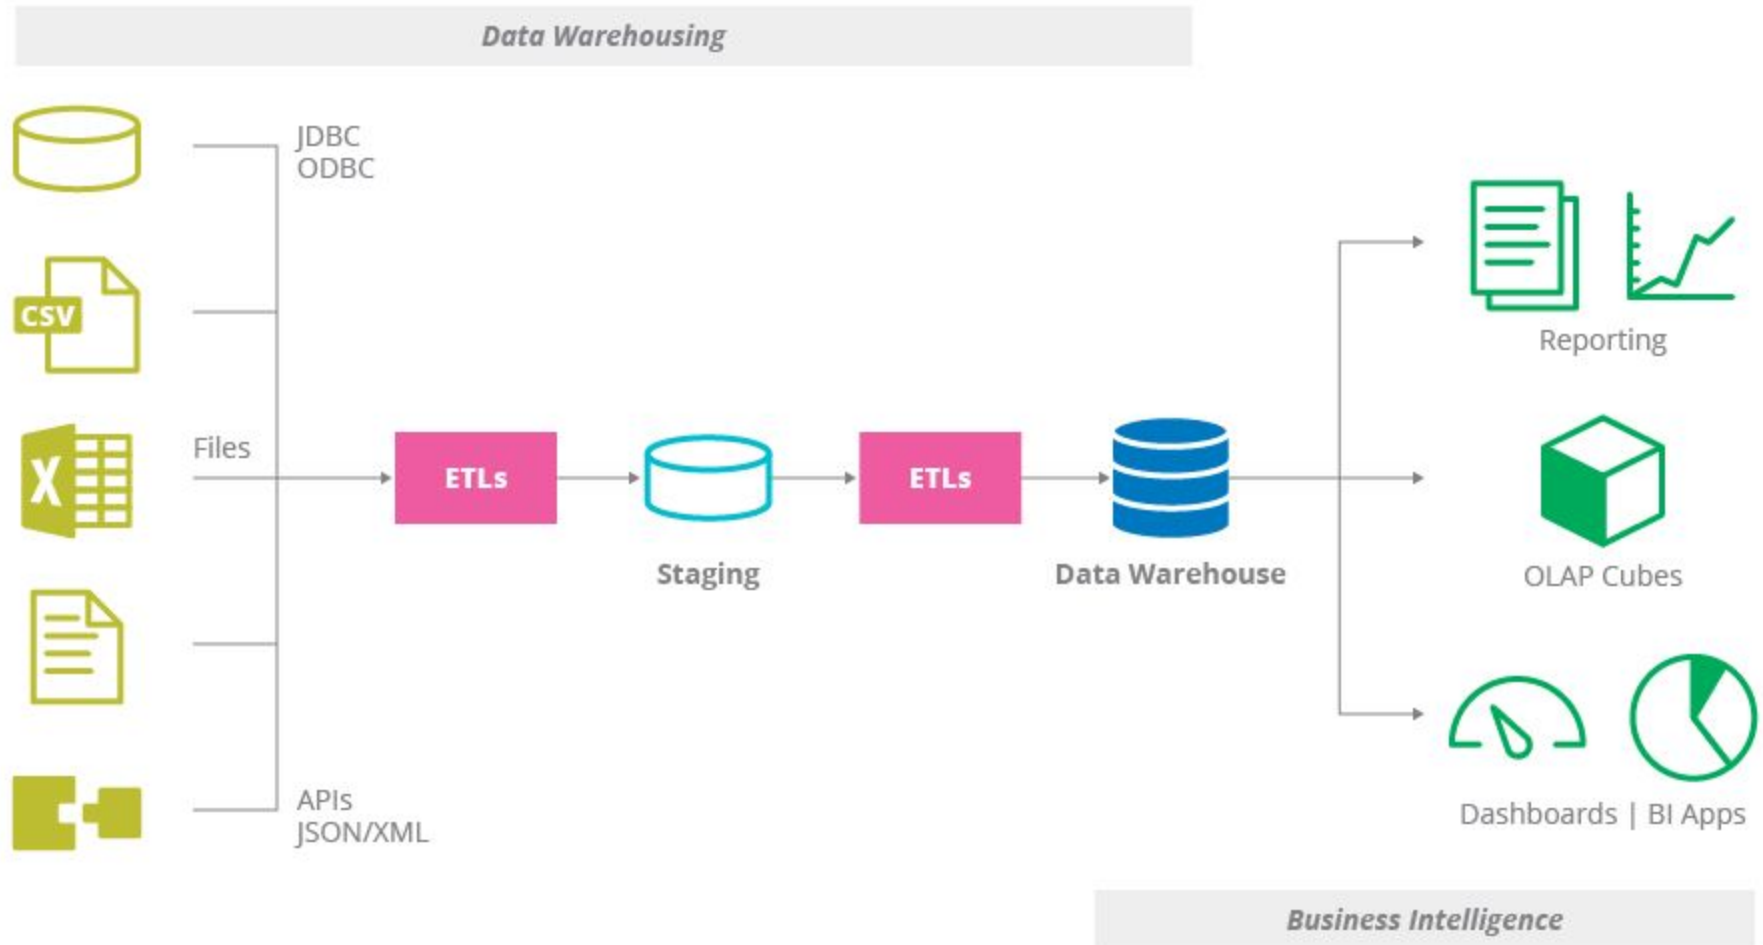
# Big Data & AI Landscape

## BIG DATA &amp; AI LANDSCAPE 2018





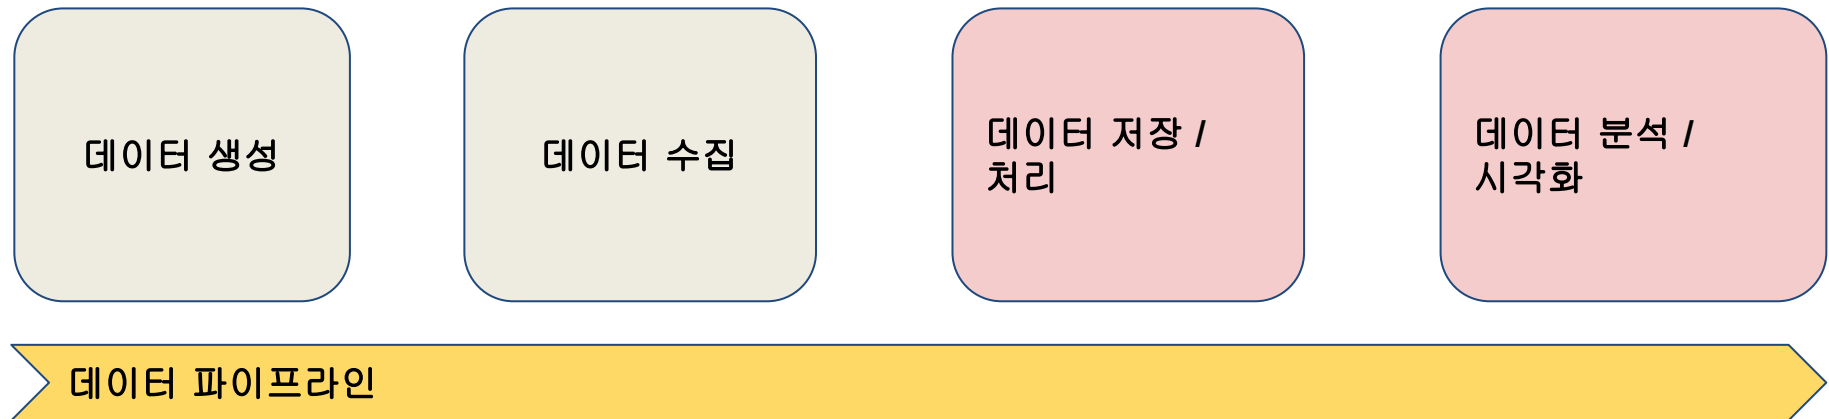
# DW & BI (Data Warehousing & Business Intelligence)



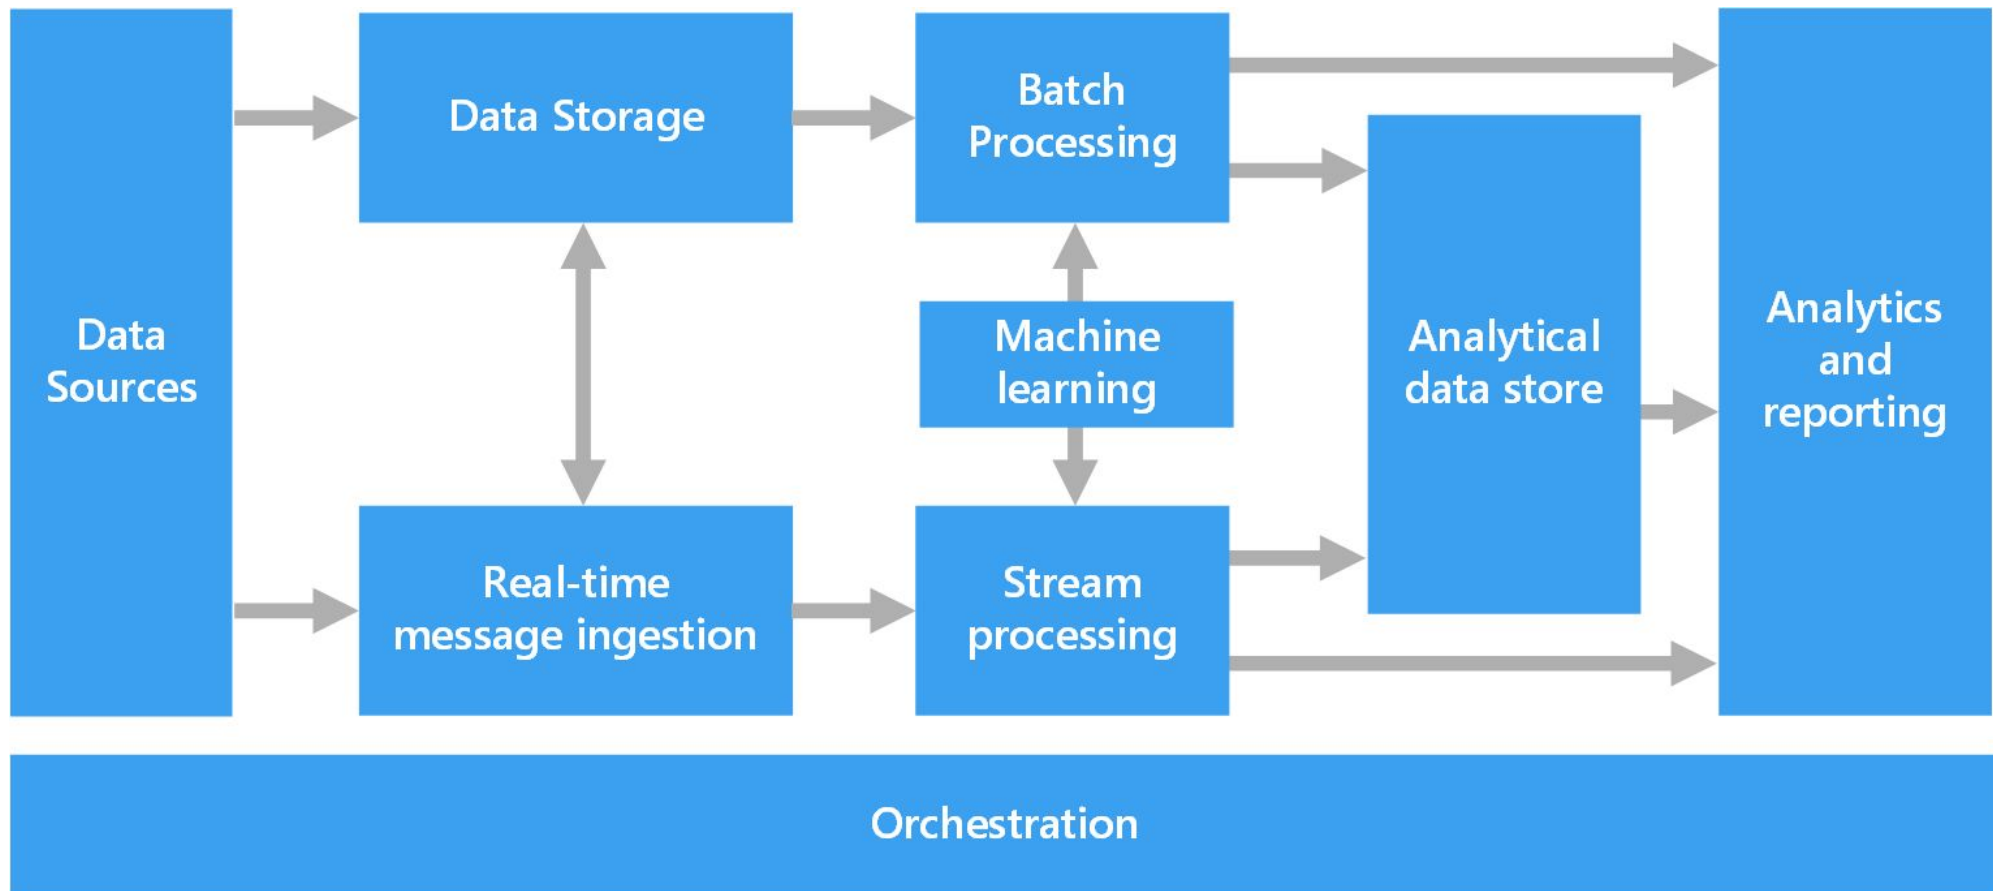
1. <https://www.thoughtworks.com/insights/blog/agile-data-warehousing-and-business-intelligence-action>

# 데이터 애플리케이션?

- 데이터 애플리케이션
  - 데이터 수집
  - 데이터 저장
  - 데이터 처리
  - 데이터 분석 및 시각화
- 빅데이터 애플리케이션(시스템) 패턴

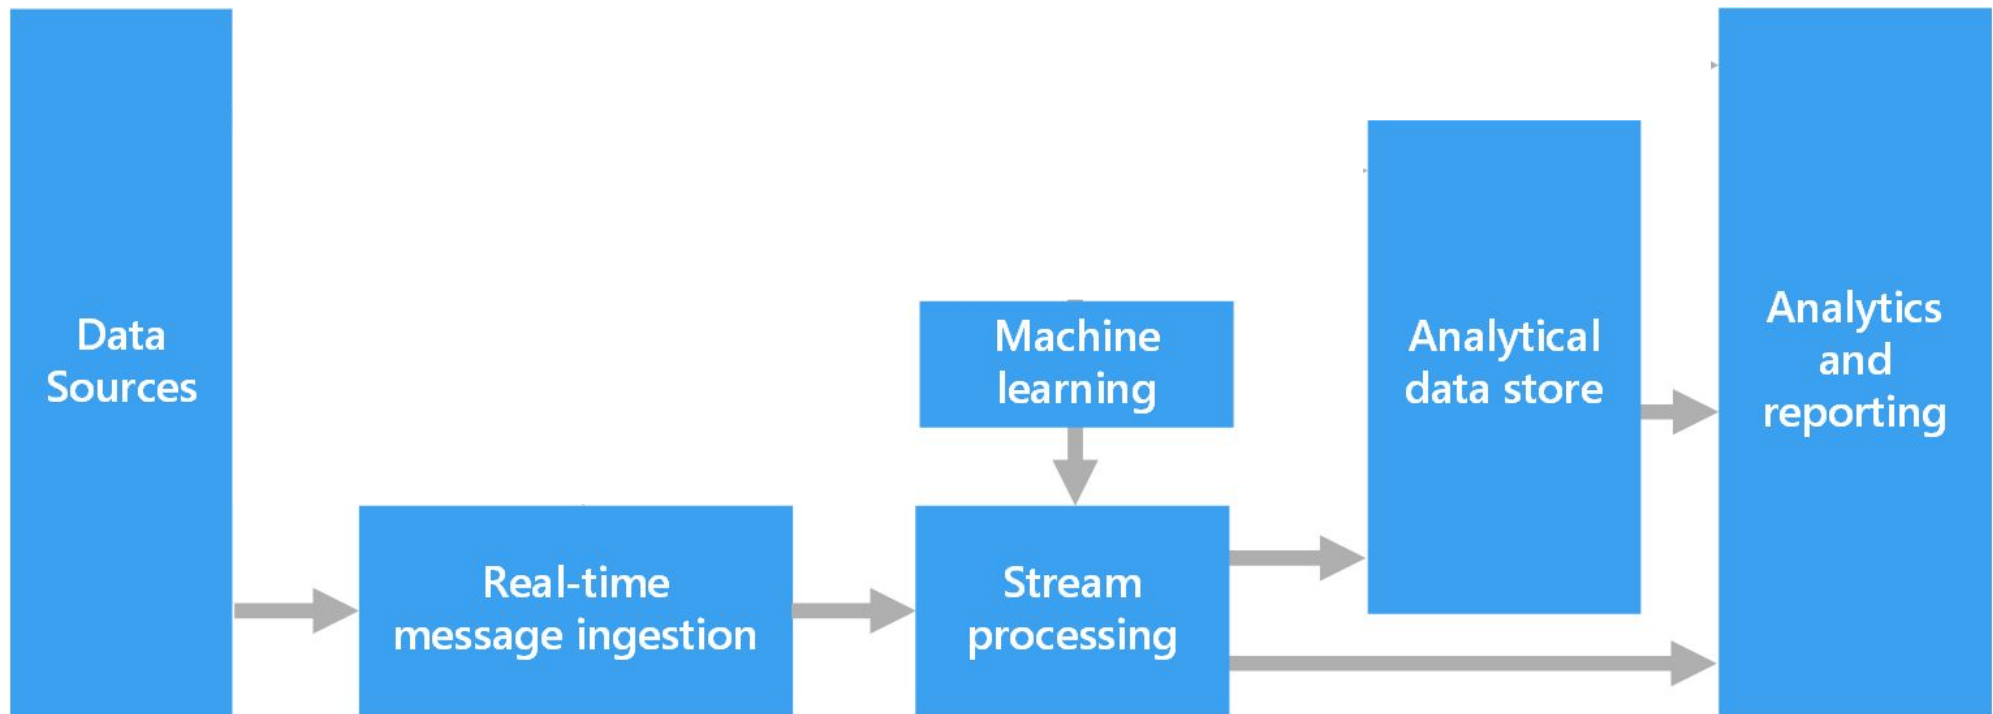


- Components of a big data architecture [1]



1. <https://docs.microsoft.com/en-us/azure/architecture/data-guide/big-data/>

- Components of the “Streaming Data” architecture





# Streaming Data Platform



# Streaming Data Platform

- Data Source
- Event Bus
  - a.k.a Message Broker
  - Transport Layer
  - Apache Kafka, Apache Pulsar, ActiveMQ, etc...
- Data Ingestion / Integration
- Data Storage
- Stream Processing
- Data Analytics / SQL / Dashboard

# Streaming Data Platform: Data Source

- **(Streaming) Data Source**

- Log files
- CDC (Transactions)
- Events (Equipments, Sensors...)
- APIs
- Agents
- Microservices
- .....

## ▪ Message Broker?, Why?

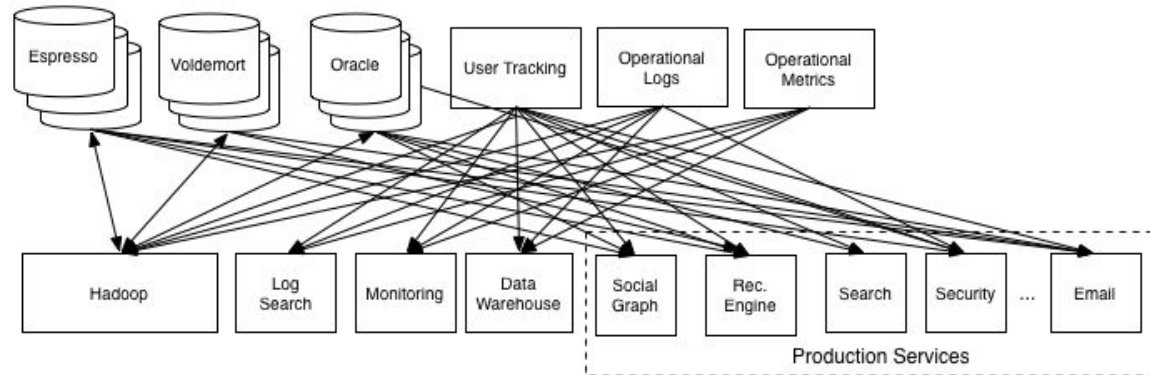
- A message broker is an architectural pattern for message validation, transformation, and routing. It **mediates communication among applications, minimizing the mutual awareness** that applications should have of each other in order to be able to exchange messages, **effectively implementing decoupling**. [1]
- Event-driven Applications
- Asynchronous Processing

1. [https://en.wikipedia.org/wiki/Message\\_broker](https://en.wikipedia.org/wiki/Message_broker)

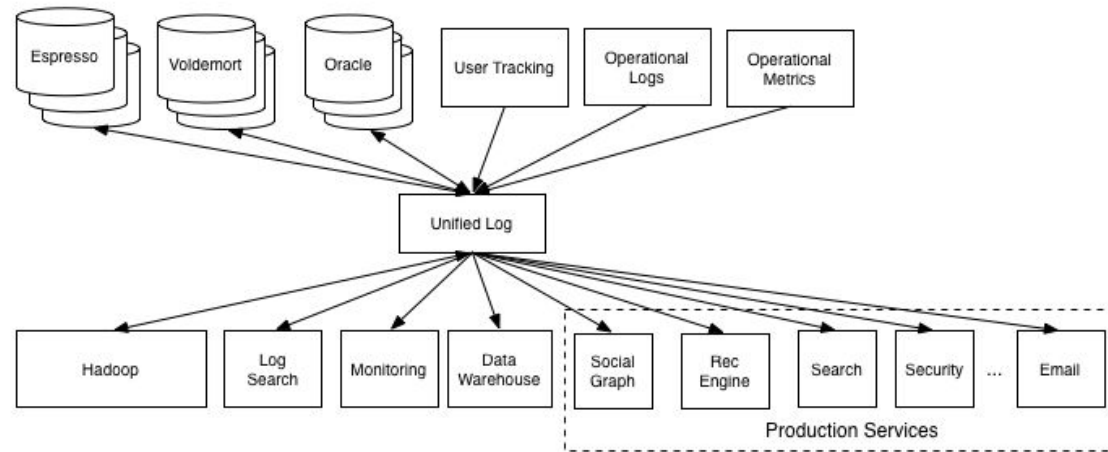


# Streaming Data Platform

Bad



Good



1. <https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>

# Streaming Data Platform

- **Message Broker**
  - a.k.a Enterprise Event Bus (Hub)
  - **Immutable append-only data store**
- **Goals for Enterprise Event Bus**
  - Buffering
  - Reliable Storage
  - Partitioning
  - Replay
  - High Throughput

# Streaming Data Platform: 이벤트 버스 = 물류창고?



- **Key properties for Next-Generation Messaging**

- High throughput
- Low latency
- Multi tenancy
- Consistency
- Ease of operations
- Resiliency
- Elasticity



- **Transmission**

- Processing Guarantees (Messaging Semantics)
- Duplication
- Latency
- Ordering

- Messaging semantics (=message delivery semantics)

## Understanding Streaming Semantics

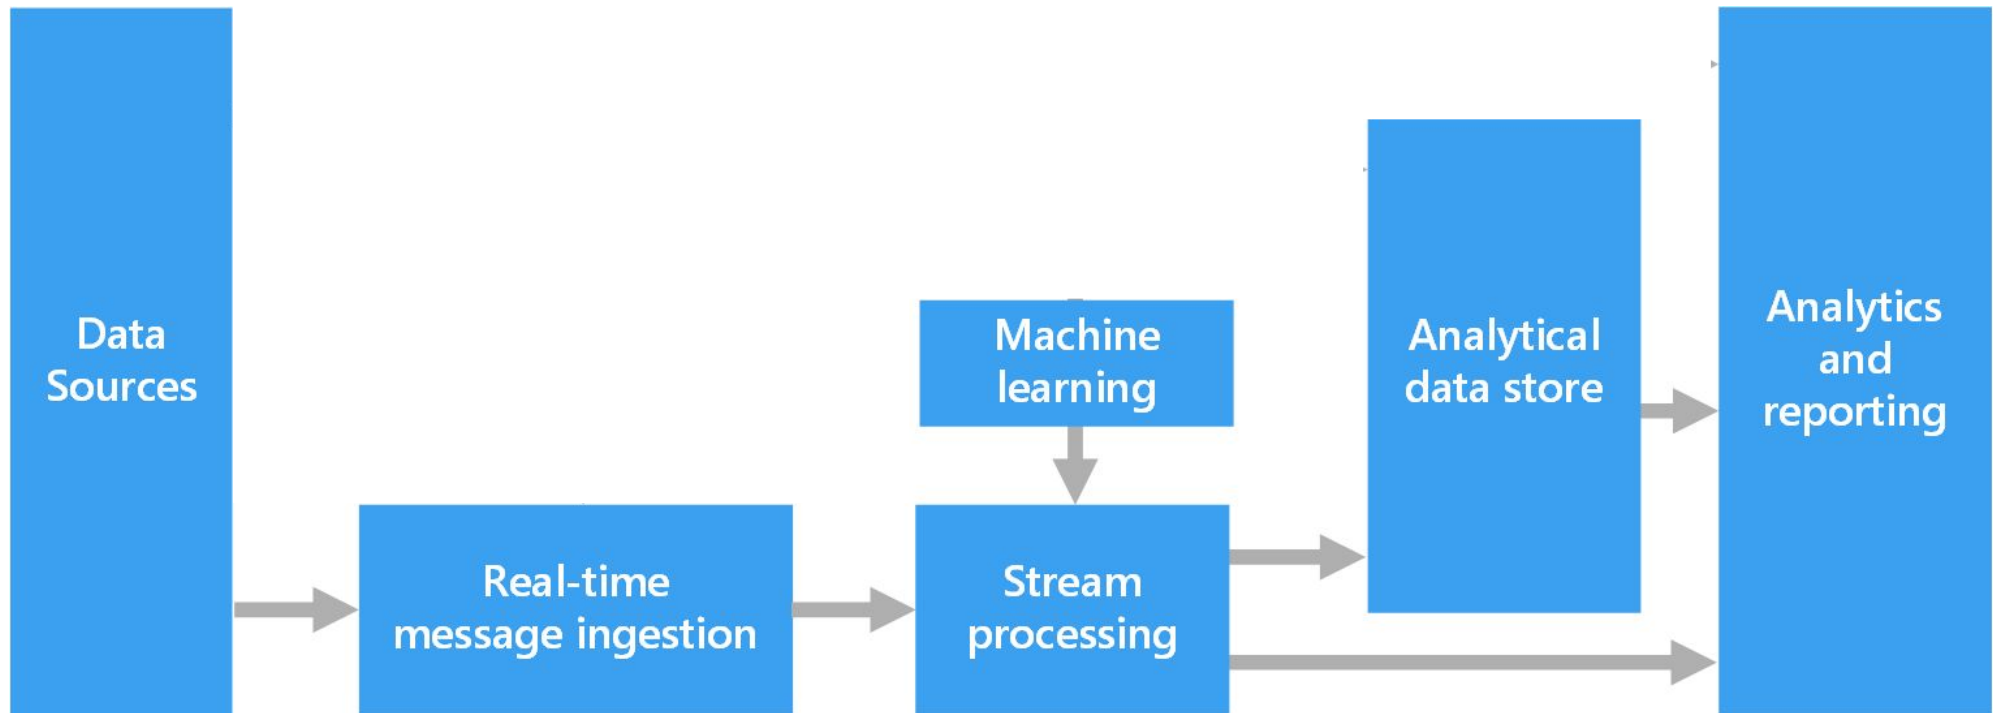


At most once	At least once	Exactly once
Message pulled once	Message pulled one or more times; processed each time	Message pulled one or more times; processed once
May or may not be received	Receipt guaranteed	Receipt guaranteed
No duplicates	Likely duplicates	No duplicates
Possible missing data	No missing data	No missing data

# OSS for Streaming Data Platform

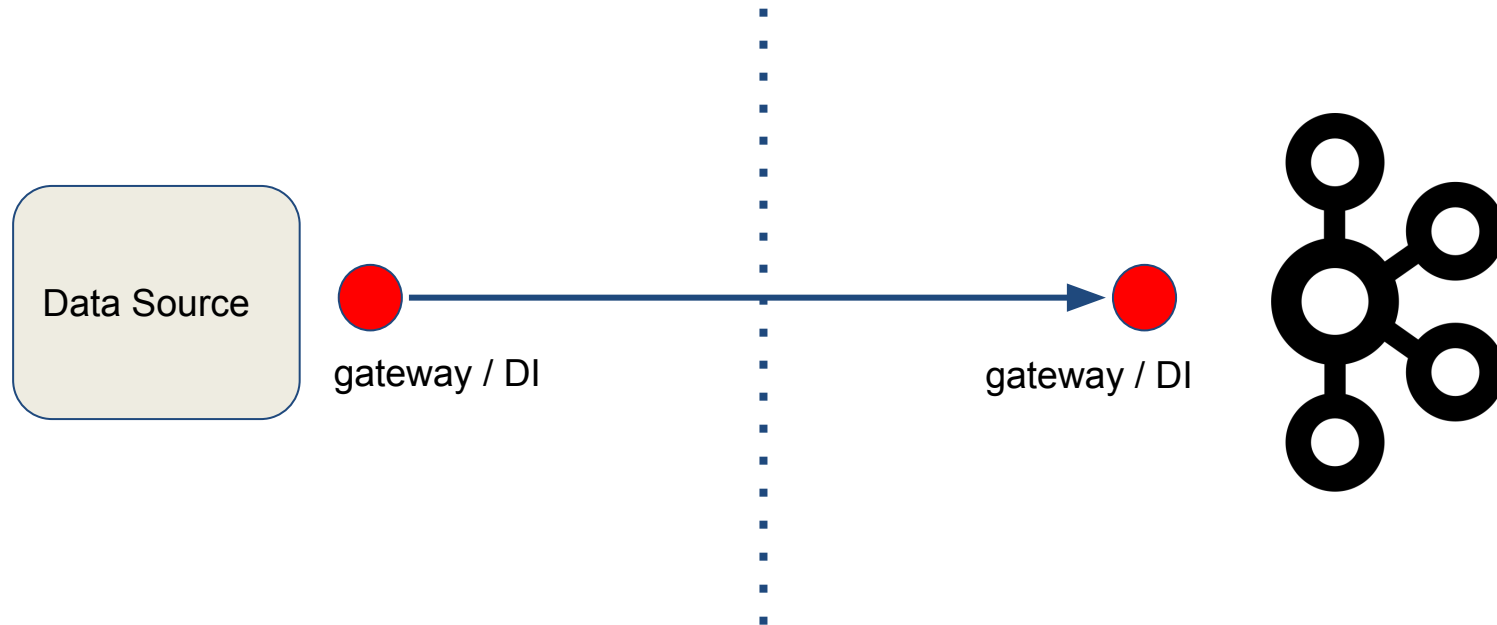


- Components of a “Streaming Data” architecture (again!)





- Data Source & Data Acquisition



- **Data Source & Data Acquisition**

- Apache Flume
- Apache Gobblin
- Kafka Connect
- DI tools...

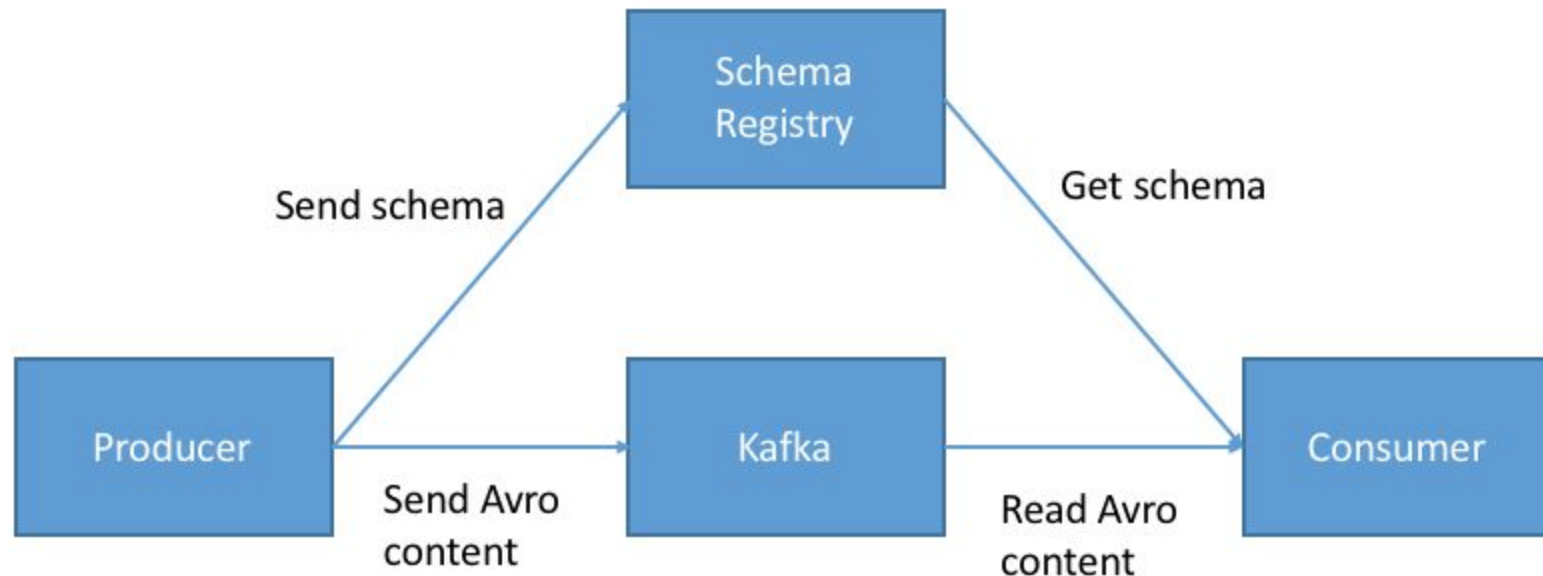
- **Message Broker / Enterprise Event Bus**

- Apache Kafka
- Apache Pulsar
- RabbitMQ
- Apache ActiveMQ
- Apache Qpid
- Apache RocketMQ
- ZeroMQ
- .....

- <https://kafka.apache.org/>
- **Kafka : High performance pub/sub messaging system**
  - Broker
  - Streams
  - Connect
- **Key terms:**
  - Message
  - Broker
  - Producer
  - Consumer
  - topic
  - topic partition
  - Consumer Group
- ***Apache Kafka vs. Confluent Platform***

# 레퍼런스 아키텍처: Apache Kafka

- **Apache Kafka + Schema Registry [1][2]**
  - Why Avro For Kafka Data? [3]



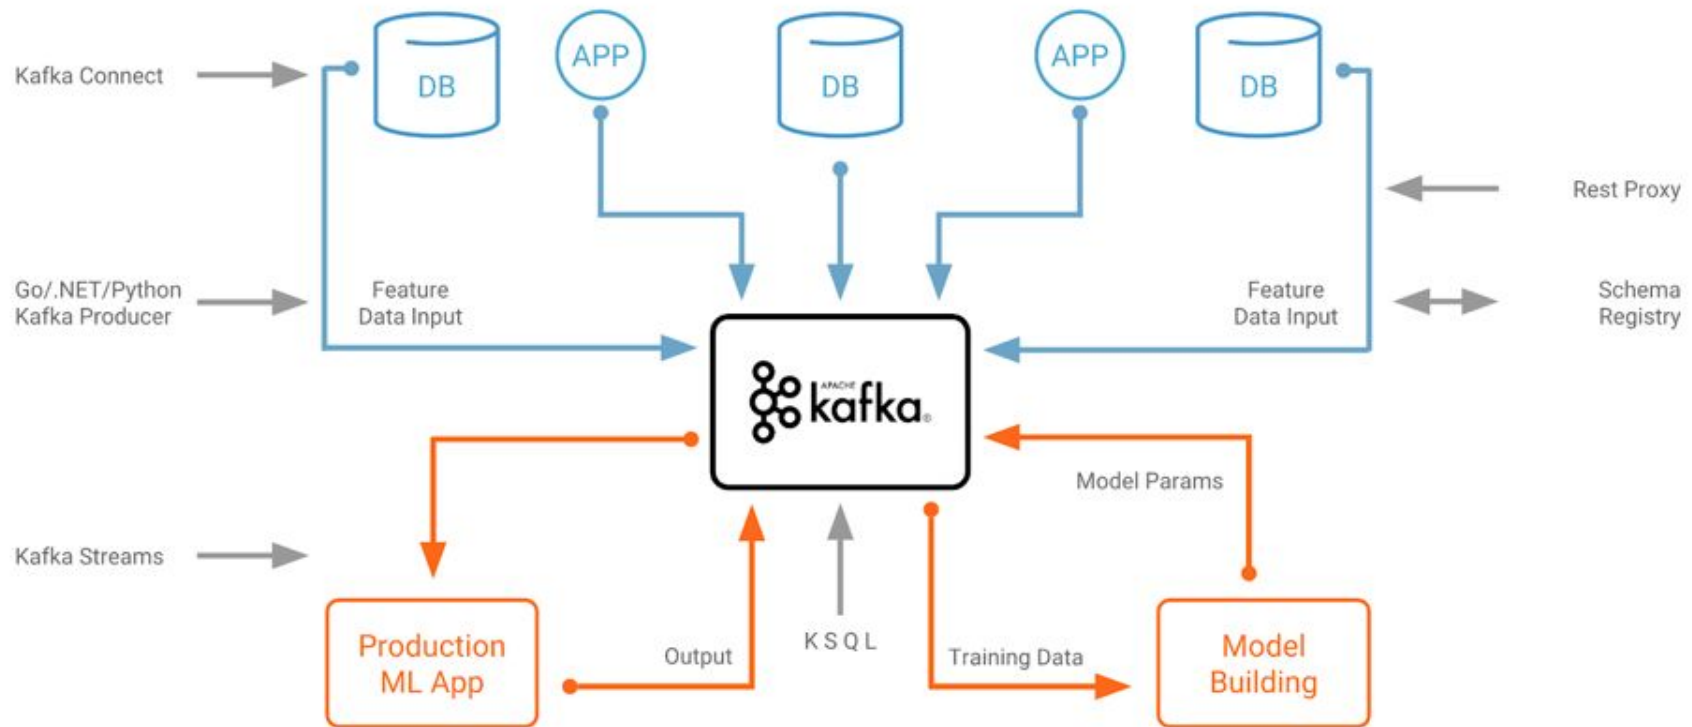
1. <https://github.com/confluentinc/schema-registry>
2. <https://medium.com/@stephane.maarek/introduction-to-schemas-in-apache-kafka-with-the-confluent-schema-registry-3bf55e401321>
3. <https://www.confluent.io/blog/avro-kafka-data/>

# 레퍼런스 아키텍처: Apache Kafka

- **Binary format for event data (Avro, Protobuf)**
  - Fast serialization / deserialization
  - more compact messages
  - have a schema of the data
  - schema evolution
  - .....
  - **Meta data and Data Discovery**



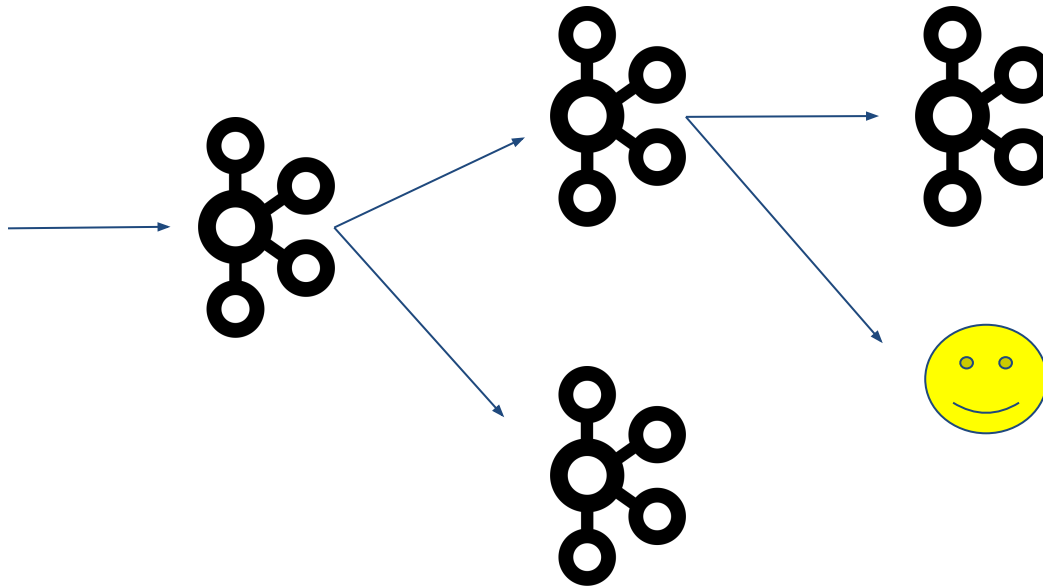
# 레퍼런스 아키텍처: Apache Kafka (for ML)



1. <https://www.confluent.io/blog/using-apache-kafka-drive-cutting-edge-machine-learning>

# Scaling Apache Kafka

- **Tuning Kafka Broker, Producer and Consumer**
  - Async producer, acks, partitions, compression, Disk I/O...
- **Tiered Kafka Cluster (=Multi-tier Architecture) [1][2]**

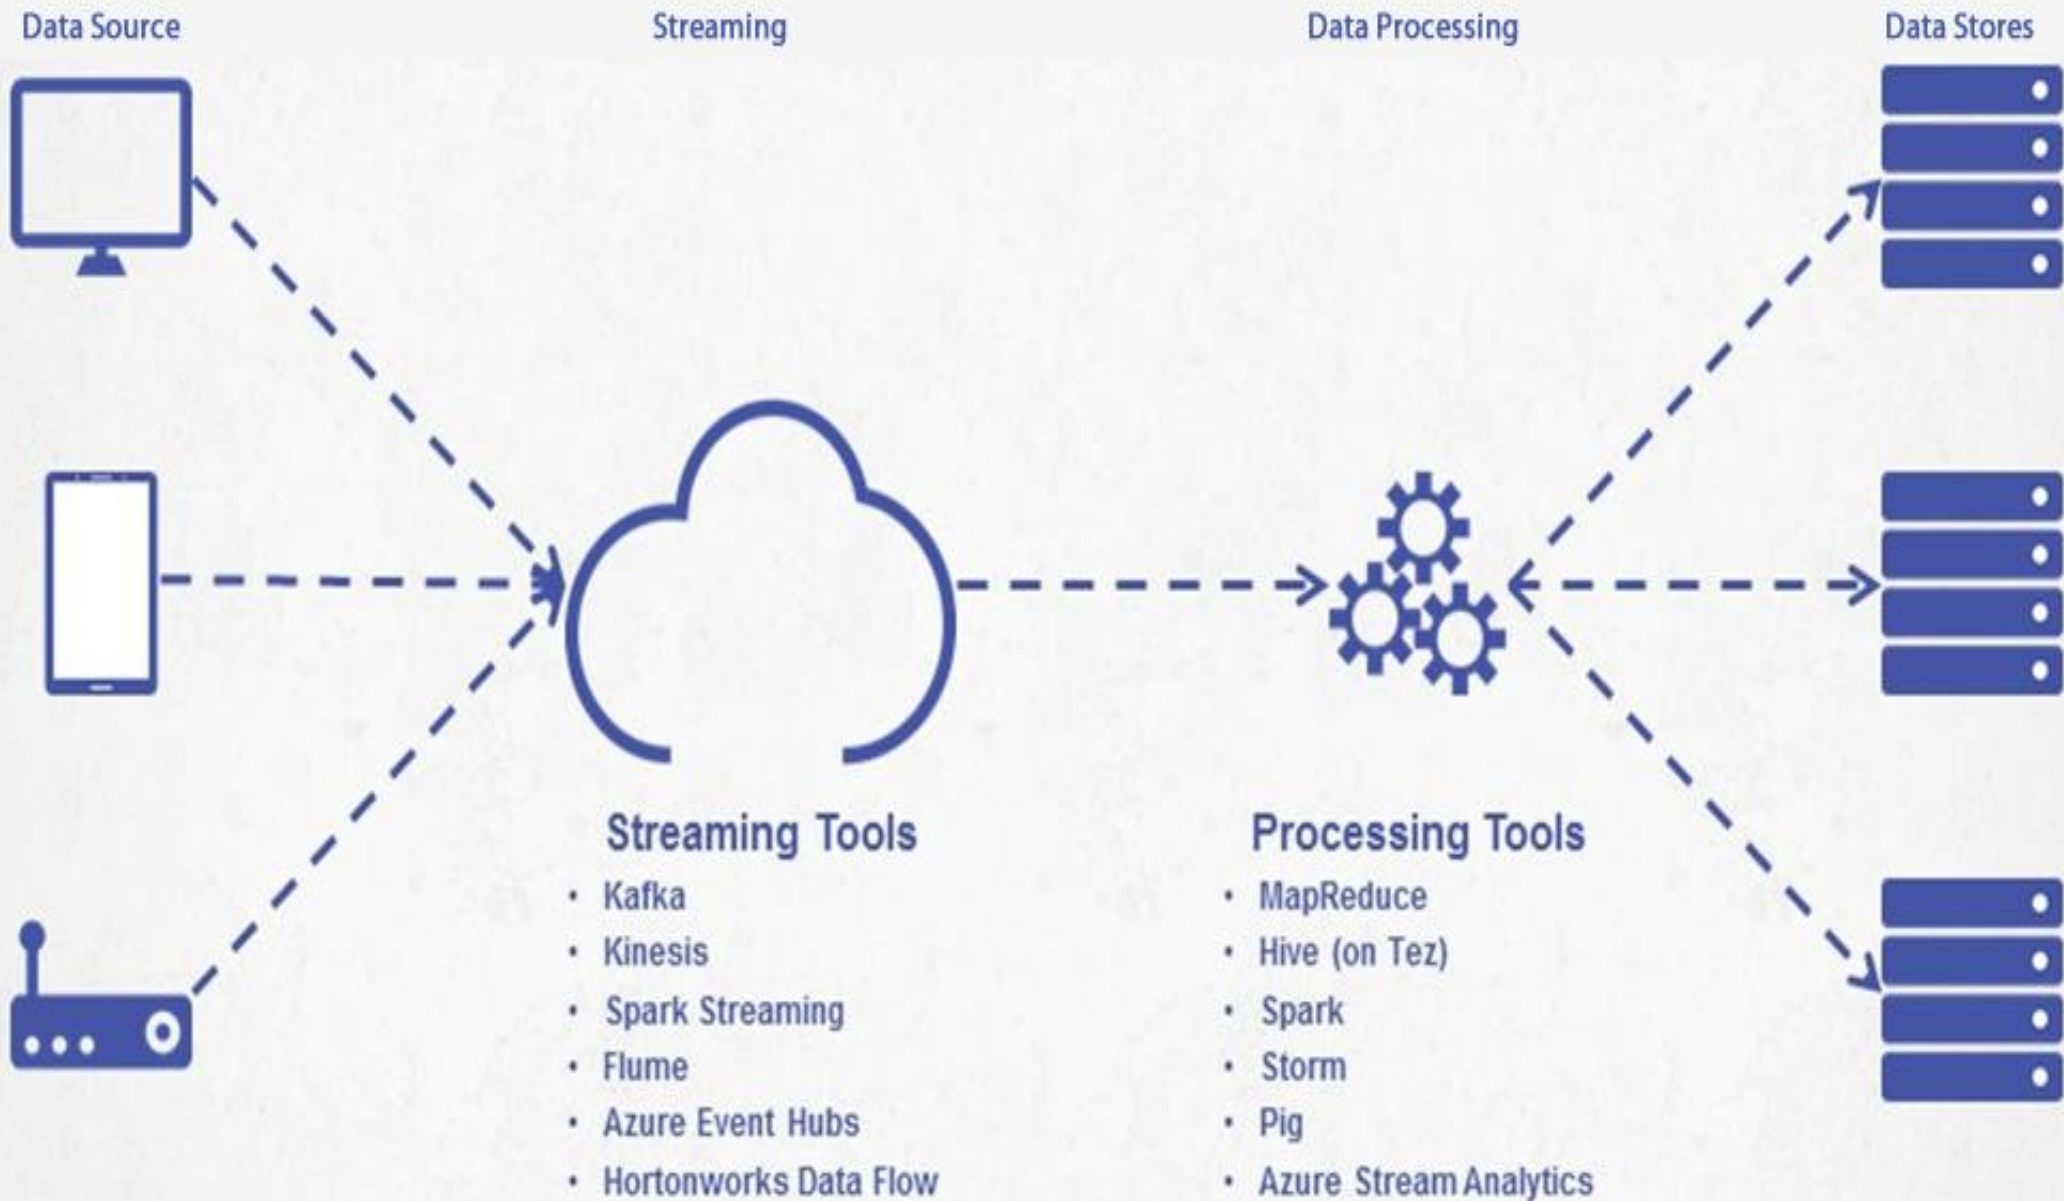


1. [Kafka at Scale](#)
2. [Building Scalable and Extendable Data Pipeline for Call of Duty Games: Lessons Learned](#)

# Streaming Data Platform: Stream Processing

- **Kafka client for Java, Scala, Python, Go, etc...**
- **Apache Spark**
- **Apache Flink**
- **Apache Samza**
- **Apache Storm**
- **Apache Heron**
- **Apache Samza**
- **Apache Beam**
- **.....**

# Streaming Data Platform: Data Store



# Streaming Data Platform: Data Store

- One size fits all?

RDBMS

DFS  
(HDFS)

Object Storage

Key/Value  
Search  
TSDB

Message  
Broker

# Streaming Data Platform: Data Store

- **RDBMS**
- **Distributed File System**
  - Hadoop HDFS
- **Object Storage**
  - S3
  - Minio
  - Ceph
- **KV Store**
  - Apache HBase
  - Apache Cassandra
  - Redis
- **Search**
  - Apache Solr
  - ElasticSearch
- **Message Broker**
  - Apache Kafka, Apache Pulsar or MQs...



# Streaming Data Platform: File Format

- **File format for Data Store**

- Distributed File System. -- HDFS, QFS, & etc...
- Object Storage -- S3 or S3 compatible data store

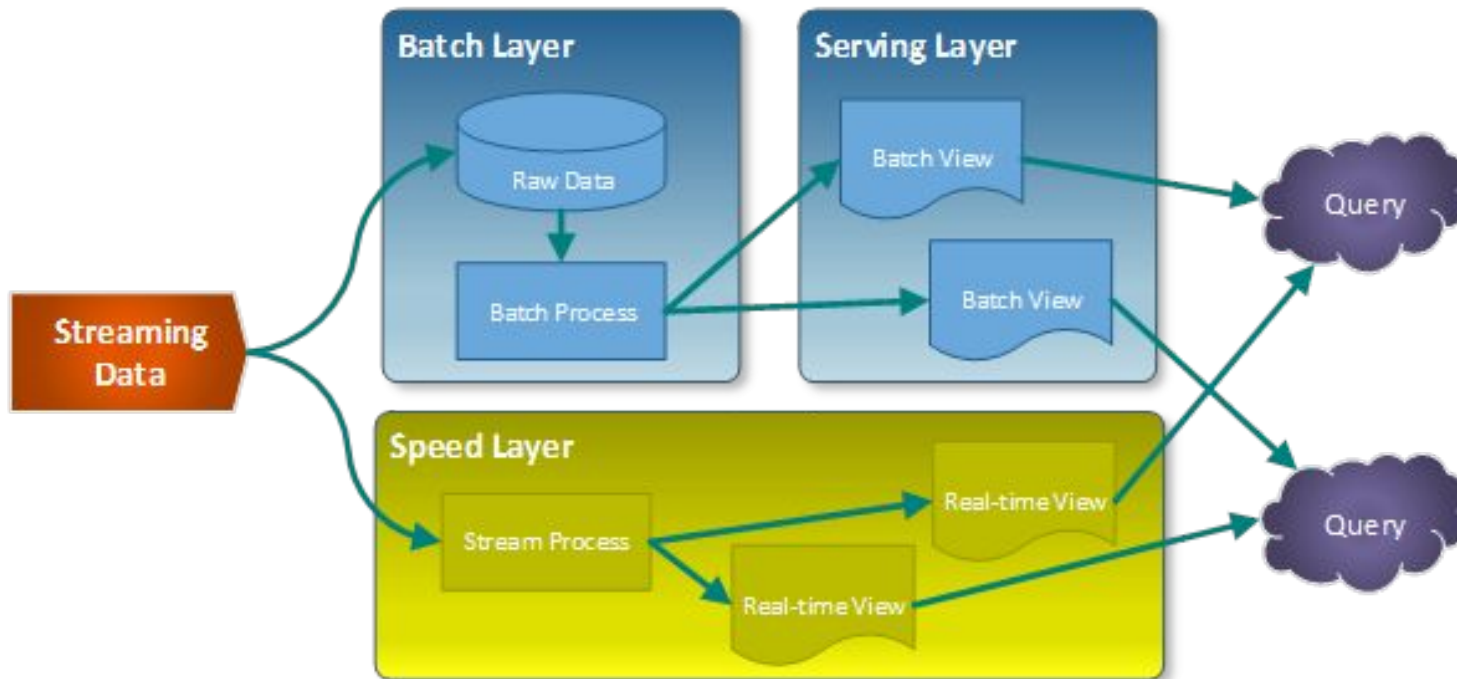
- **Big Data File Formats**

- Text (CSV, JSON)
- Apache Avro
- Apache ORC\*\*
- Apache Parquet\*\*

**\*\* Columnar storage > immutable data and analytics queries**

# 레퍼런스 아키텍처: Lambda Arch. vs Kappa Arch.

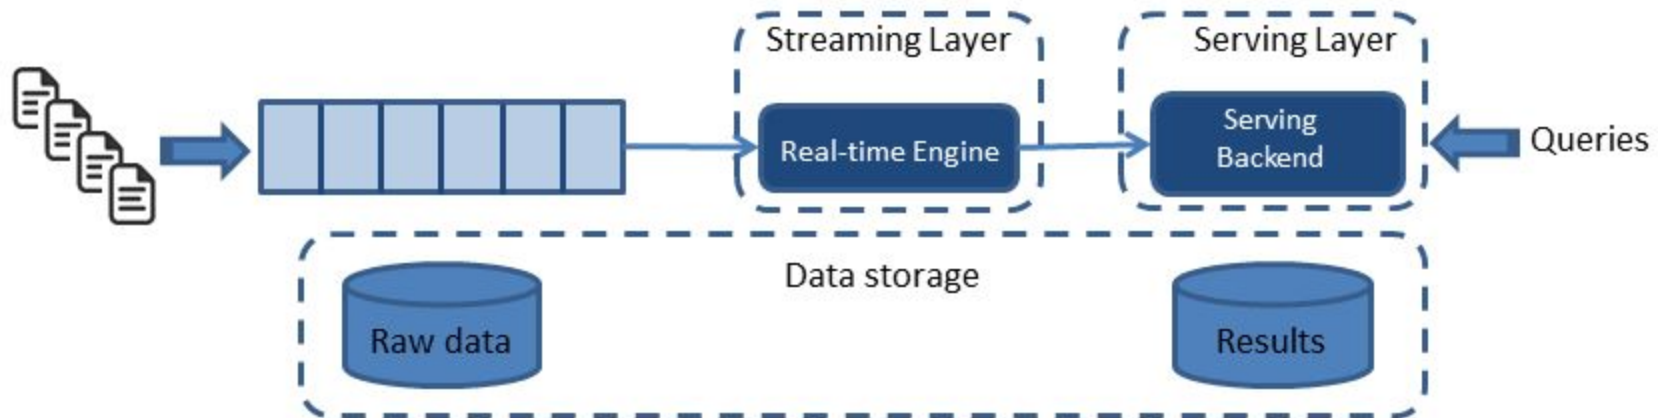
- Lambda Architecture



<https://www.talend.com/blog/2017/08/28/lambda-kappa-real-time-big-data-architectures/>

# 레퍼런스 아키텍처: Lambda Arch. vs Kappa Arch.

## ▪ Kappa Architecture



1. <https://www.oreilly.com/ideas/applying-the-kappa-architecture-in-the-telco-industry>

- **Real-time Data Ingestion / Integration**

- Apache Flume
- Apache Gobblin
- Kafka Connect
- StreamSets
- .....

# Streaming Data Platform: Data Apps (Analytics / SQL / Dashboard)

- **Data Applications**
  - API
- **Querying Data**
  - SQL
- **Jupyter notebooks**
- **Data Discovery**
  - [metatron discovery](#)
- **ML / DL**
- **Search**
- **Dashboard**

# Stream analytics @ Public Cloud

- **AWS**

- <https://aws.amazon.com/ko/streaming-data/>

- **MS Azure**

- <https://azure.microsoft.com/en-in/services/stream-analytics/>

- **GCP**

- <https://cloud.google.com/solutions/big-data/stream-analytics/>



# Hands-on Labs



- **Stock Trading(Quote) API를 이용한 실시간 데이터 수집 및 분석**

- Data Sources

- Dataset
- API

- Real-time message ingestion

- Stream Ingestion

- Kafka Clients

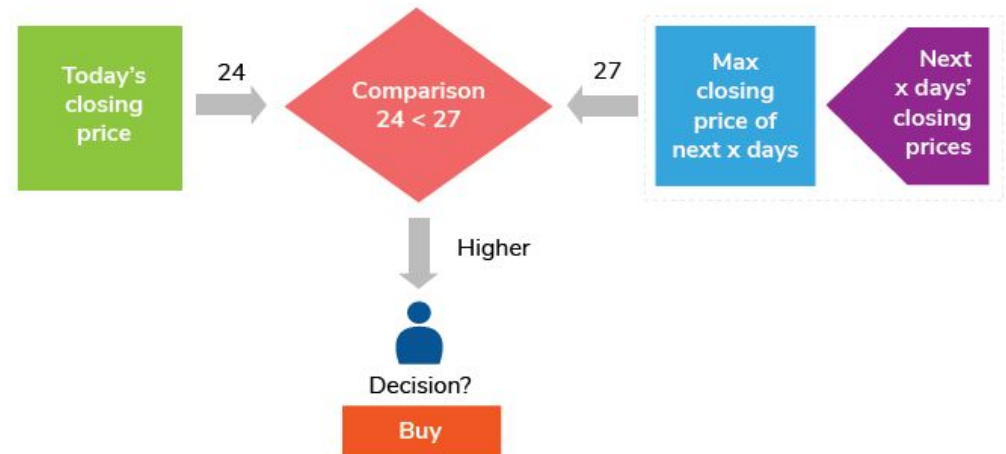
- Event Bus

- Apache Kafka



## ▪ Stock Trading

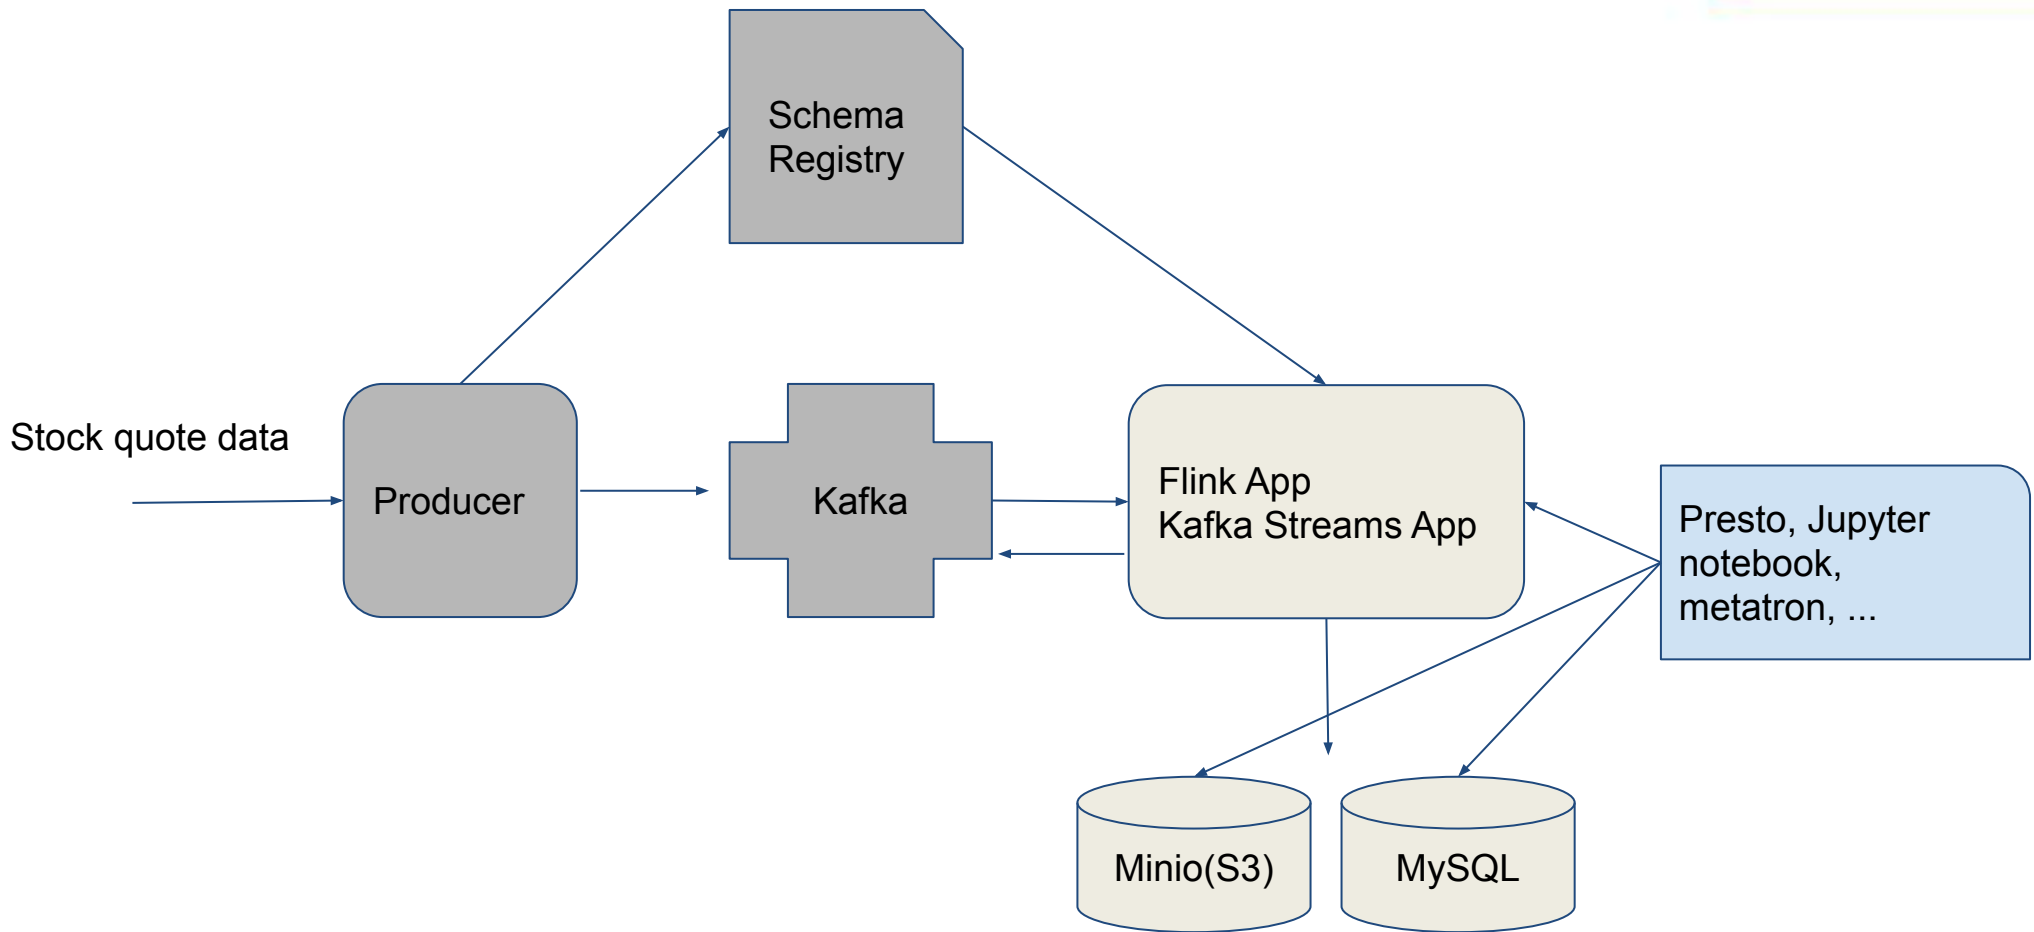
- Stream processing
  - Kafka Streams, Apache Spark, Apache Flink
- Data Store
  - RDBMS, Object Storage, RDBMS
- Data Analytics
  - SQL
  - Python



# Hands-on Labs

- <https://github.com/youngwookim/skbsdc-bda-2019>
  - \$ cd /path/to/workspace
  - \$ git checkout https://github.com/youngwookim/skbsdc-bda-2019.git

# Hands-on Labs: Data flow

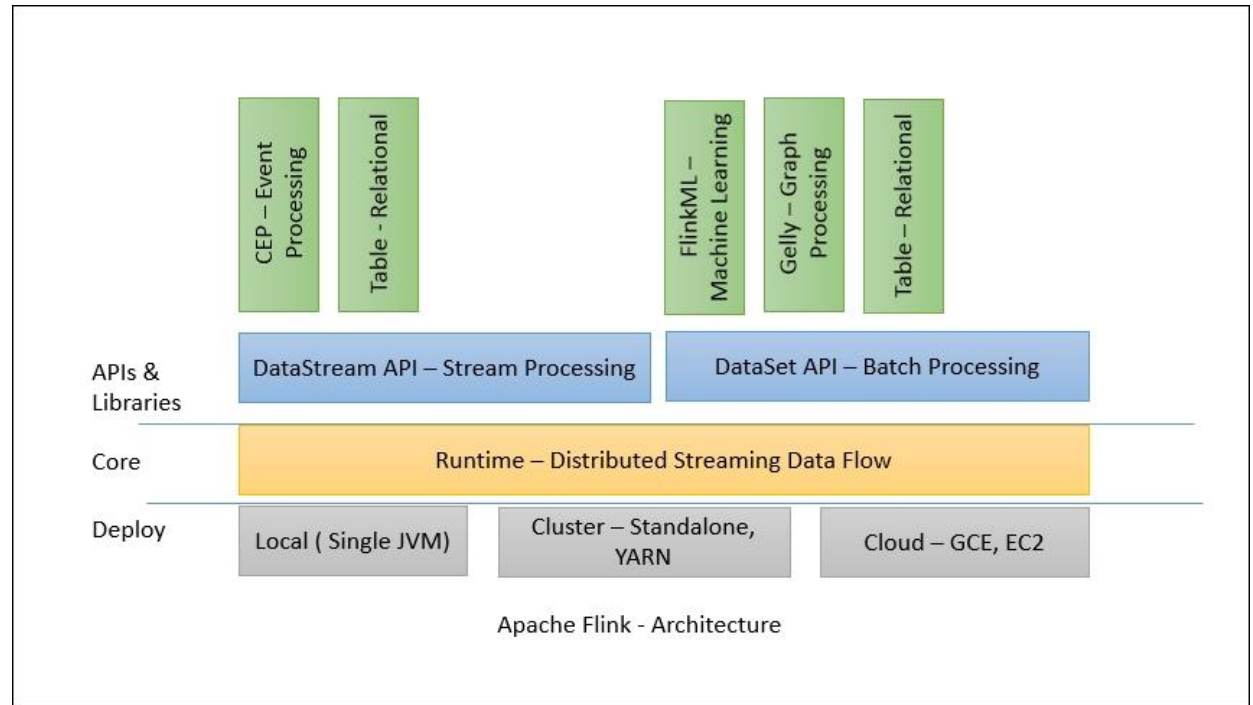






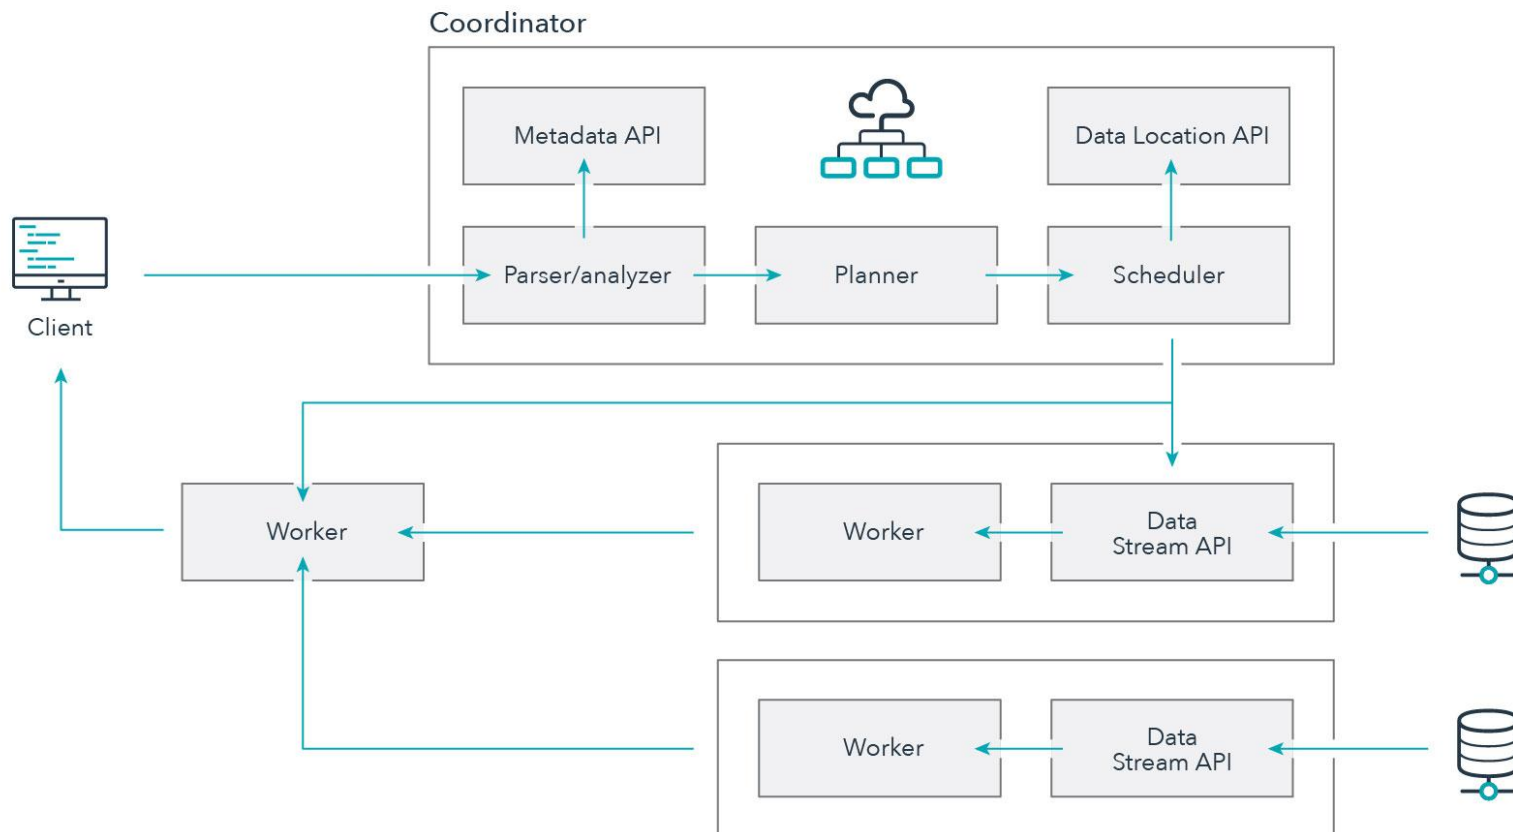
# Apache Flink

- <https://flink.apache.org/>
  - vs. Apache Spark
  - vs. Apache Samza
  - vs. Apache Storm
  - vs. Apache Heron



# Presto

- **a.k.a PrestoSQL**
  - <http://prestosql.io>



# Metatron Discovery

- **metatron discovery**

- <https://www.youtube.com/channel/UC5ldHK8qBiN9zVgD7SvO41g>

## Refs.

1. [Big data architectures](#)
2. [Architecting a next-generation data platform](#)
3. [Designing modern streaming data applications](#)
4. [Kafka Best Practices](#)
5. [Kafka Best Practices](#)
6. [I Heart Logs](#)
7. <https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>
8. [Architectural considerations for Hadoop Applications](#)

# Image Credits

1. <http://insight360.com/big-data/data-stores/>
2. [https://en.wikipedia.org/wiki/Apache\\_Kafka](https://en.wikipedia.org/wiki/Apache_Kafka)
3. <https://www.talend.com/blog/2017/08/28/lambda-kappa-real-time-big-data-architectures>
4. [https://subscription.packtpub.com/book/big\\_data\\_and\\_business\\_intelligence/9781786466228/1/ch01lv1sec8/architecture](https://subscription.packtpub.com/book/big_data_and_business_intelligence/9781786466228/1/ch01lv1sec8/architecture)
5. <https://www.starburstdata.com/learn-presto/reference-architectures/>